



# **3547 Intelligent Agents & Reinforcement Learning**

## **Module 4: Planning and Knowledge Representation**



# Course Plan

## Module Titles

Module 1 – Introduction to Intelligent Agents

Module 2 – Current Focus: Search

Module 3 – Logical Inference

**Module 4 – Planning and Knowledge Representation**

Module 5 – Probabilistic Reasoning

Module 6 – Intro to Reinforcement Learning and Finite Markov Decision Processes

Module 7 – Dynamic Programming and Monte Carlo Methods

Module 8 – Temporal Difference Learning

Module 9 – Function Approximation for RL

Module 10 – Deep Reinforcement Learning and Policy Gradient Methods

Module 11 – Introduction to Advanced DRL

Module 12 – Presentations (no content)



# Learning Outcomes for this Module

- Introduce you to expert systems
- Extend your skills working with Propositional Logic to First Order Logic (FOL)
- Introduce Prolog
- Enable you to model the passage of time in logic
- Use logic programming to plan sequences of action for an agent
- Show you how to model knowledge using ontologies and relationships



# Topics for this Module

- **4.1** Expert Systems and First Order Logic
- **4.2** Introduction to Prolog
- **4.3** Modeling the Flow of Time
- **4.4** Planning
- **4.5** Knowledge Representation
- **4.6** Resources and Wrap-up



## Module 4 – Section 1

# Expert Systems and First Order Logic

# Expert Systems

- Need to be able to deal with uncertainty and be transparent
- Components
  - Problem Solving: DB + inference engine
  - User Interaction: UI + query capability
- If-then rules are a natural choice
- MYCIN
- AL3
- Knowledge acquisition is the bottleneck

# Terminology for First Order Logic

- **Constant:** A symbol that stands for an object
- **Variable:** A symbol representing a variable term (and in FOL can serve as an argument to a function)
- **Predicate:** A symbol which stands for a relation of form  $p(t_1 \dots t_n)$  where  $p$  is a predicate symbol and each  $t_i$  is a term (called an **argument** to the predicate)
- **Function:** A symbol that stands for a function of form  $f(t_1 \dots t_n)$  and represents an object
- **Term:** A constant, a variable or a function; a logical expression that refers to an object
- **Atom (or atomic symbol):** a formula of form  $p$
- **Expression:** Either a term, an atom, a definite clause, or a query
- **Ground:** An expression is ground if doesn't contain any variables
- **Interpretation:** Specifies which objects, relations, and functions are referred to by the constant, predicate and function symbols
- The definitions of definite clause, knowledge base, rule and query remain the same as for propositional calculus



## Module 4 – Section 2

# Introduction to Prolog

# Prolog Hands-On

- <https://swish.swi-prolog.org/>

a :- b, c.

b :- d, e.

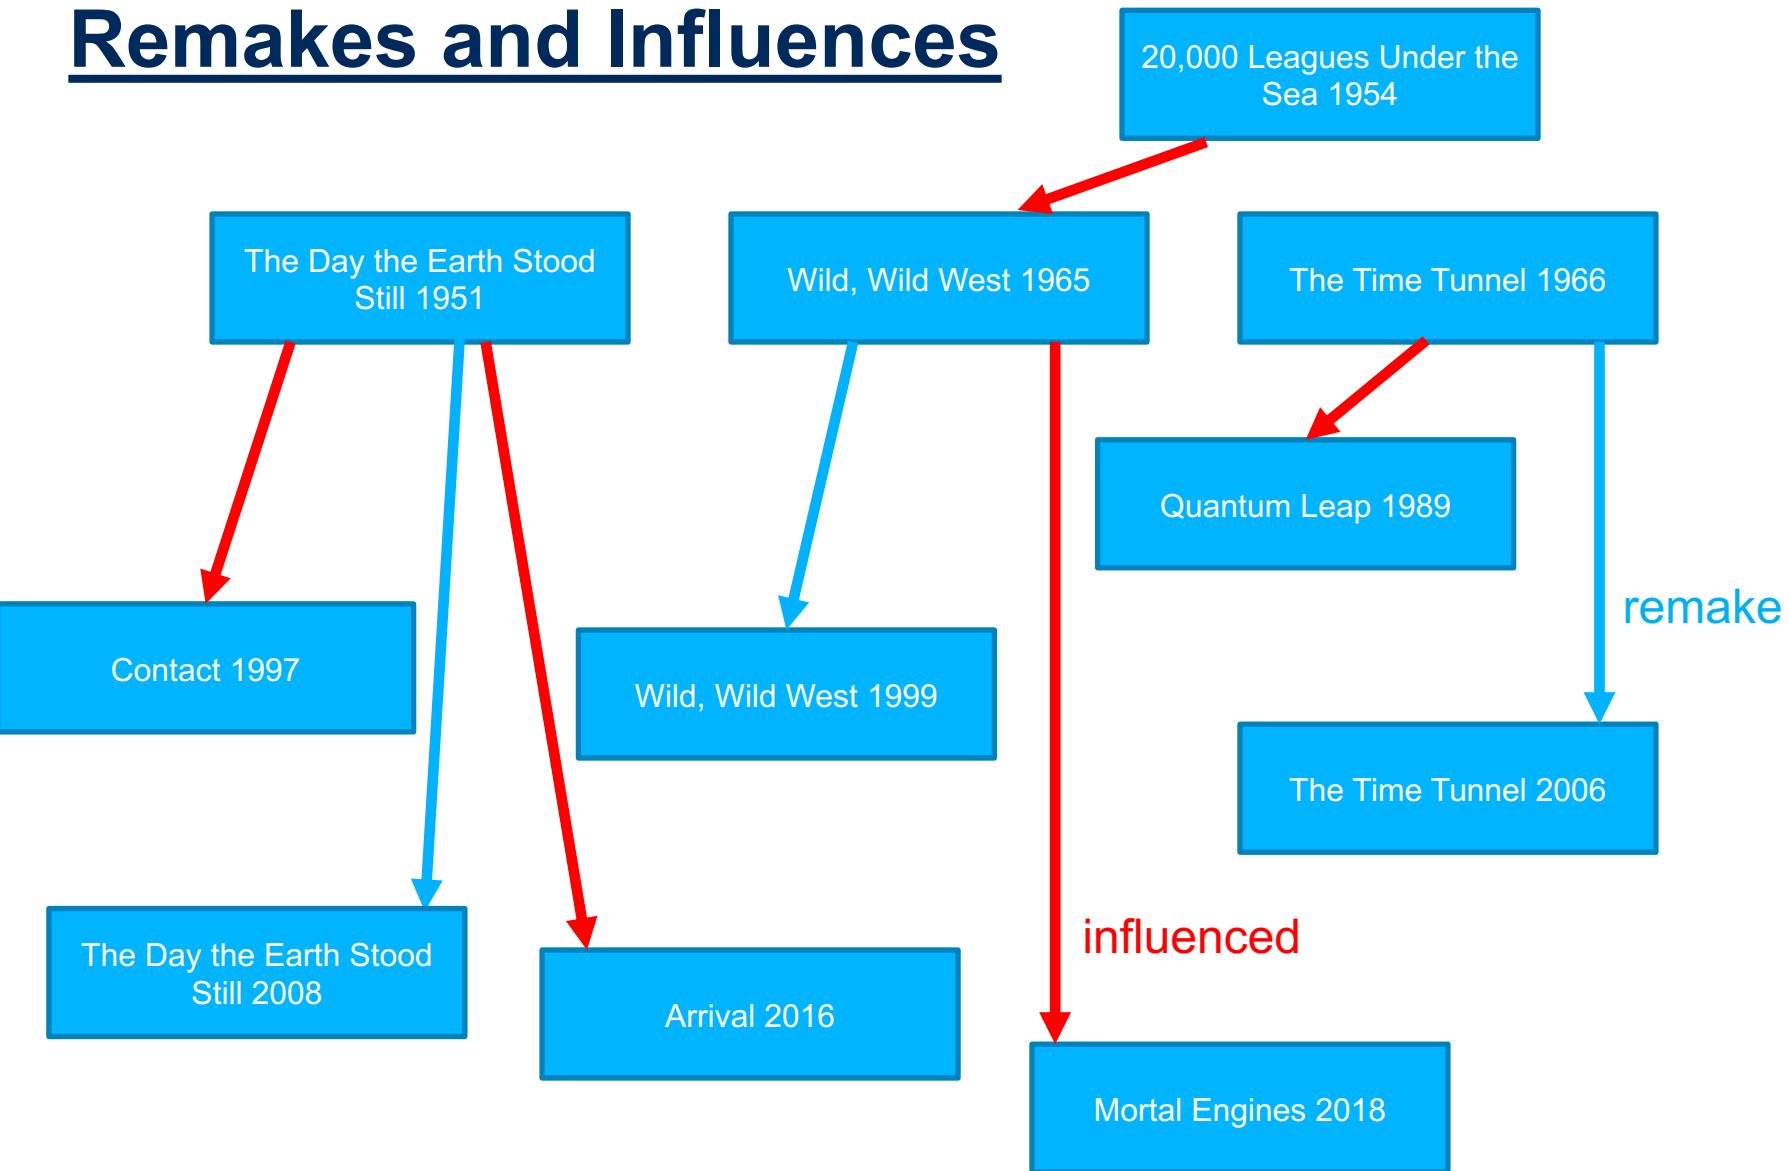
c :- e.

d.

e.

f :- a, g.

# Remakes and Influences



All opinions on influences are the slide author's

# Prolog Relations

- <https://swish.swi-prolog.org/>

remade(day\_earth\_51, day\_earth\_08).

remade(wild\_west\_65, wild\_west\_99).

remade(time\_tunnel\_66, time\_tunnel\_06).

influenced(leagues\_54, wild\_west\_65).

influenced(day\_earth\_51, arrival\_16).

influenced(wild\_west\_65, mortal\_engines\_18).

influenced(time\_tunnel\_66, quantum\_leap\_89).

influenced(day\_earth\_51, contact\_97).

# Prolog Queries

?-remade(wild\_west\_65, wild\_west\_99).

?-remade(wild\_west\_65, quantum\_leap\_89).

?-directed(wild\_west\_65, quantum\_leap\_89).

?-influenced(leagues\_54, arrival\_16).

?-influenced(X, mortal\_engines\_18).

?-influenced(X, mortal\_engines\_18), influenced(A, X).

Now: Formulate a query to answer:

- Which movie influenced both Arrival and Contact?

# Facts and Rules

excellent(arrival\_16).

- How would you express, (using :-)  
For all X and Y,  
    X is excellent if  
        X influenced Y, and Y was excellent.
- How would you ask what movies were also influenced by the movies that influenced Arrival?

# Recursive Rules

`led_to(X,Y):-influenced(X,Y).`

`led_to(X,Y):-influenced(X,A), influenced(A,Y).`

`led_to(X,Y):-influenced(X,A), influenced(A,B), influenced(B,Y).`

`led_to(X,Y):-influenced(X,A), influenced(A,B), influenced(B,C),  
influenced(C,Y).`

...

Define with two rules:

`led_to(X,Y):-influenced(X,Y).`

`led_to(X,Y):-influenced(X,A), led_to(A,Y).`

# Try This:

```
word(d,o,g).  
word(r,u,n).  
word(t,o,p).  
word(f,i,v,e).  
word(f,o,u,r).  
word(l,o,s,t).  
word(m,e,s,s).  
word(u,n,i,t).  
word(b,a,k,e,r).  
word(f,o,r,u,m).  
word(g,r,e,e,n).  
word(s,u,p,e,r).  
word(p,r,o,l,o,g).  
word(v,a,n,i,s,h).  
word(w,o,n,d,e,r).  
word(y,e,l,l,o,w).
```

```
solution(L1,L2,L3,L4,L5,L6,L7,L8,L9,L10,L11,L12,L13,L14,L15,L16) :-  
    word(L1,L2,L3,L4,L5),  
    word(L9,L10,L11,L12,L13,L14),  
    word(L1,L6,L9,L15),  
    word(L3,L7,L11),  
    word(L5,L8,L13,L16).  
?- solution(L1,L2,L3,L4,L5,L6,L7,L8,L9,L10,L11,L12,L13,L14,L15,L16).
```

L1	L2	L3	L4	L5	
L6		L7		L8	
L9	L10	L11	L12	L13	L14
L15				L16	

# Declarative and Procedural Meaning

- Consider  $X :- A, B$  where  $X, A, B$  are terms
- We can read it declaratively:
  - if  $A$  is true and  $B$  is true then  $X$  is true
  - $A$  and  $B$  imply  $X$
- Or procedurally:
  - To solve problem  $X$ , *first* solve subproblem  $A$  *then* subproblem  $B$
  - To satisfy  $X$ , *first* satisfy  $A$  *then*  $B$

# Disjunction in Prolog

- Prolog also supports disjunction (or):
  - Use semicolon rather than colon
  - e.g.  $X :- A; B.$
  - Has the same meaning as:  
 $X :- A.$   
 $X :- B.$
- Comma binds stronger than semicolon
- e.g.  $X :- M, N; A, B, C.$  means  $X :- (M,N);(A,B,C).$

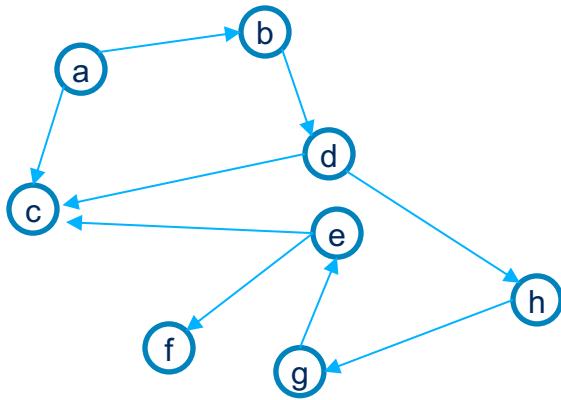
# How Prolog Scans

```
movie(oblivion, cruise).  
movie(vanilla_sky, cruise).  
movie(edge_of_tomorrow,cruise).  
movie(paycheck, affleck).  
movie(dumb_and_dumber, carrey).  
movie(shakespeare_in_love, affleck).  
nationality(affleck, american).  
nationality(cruise, american).  
nationality(carrey, canadian).  
  
movie(X,Y),nationality(Y,canadian).
```

# Other Prolog Features

- Lists
  - [item\_one, item\_two, item\_three]
  - Constructor is . e.g. .(item\_one, (time\_two, (item\_three, [ ])))
  - Head and tail operator is | e.g. L = [a|Tail]
  - Concatenate with conc(a, b, c).
- Arithmetic
  - Use *is* to evaluate e.g. X is 1 + 2
  - Use =:= to compare numbers by evaluating them

# Graph Search in Prolog



```
link(a,b).  
link(a,c).  
link(b,d).  
link(d,c).  
link(d,h).  
link(h,g).  
link(g,e).  
link(e,c).  
link(e,f).
```

```
path(Node, Node).
```

```
path(StartNode, EndNode) :- link(StartNode, NextNode), path(NextNode, EndNode).
```

# WumpusWorld in Prolog

<https://www.metalevel.at/wumpusworld/wumpus.pl>

<https://github.com/rlegendi/wumpus-prolog/blob/master/wumpus.pl>



UNIVERSITY OF TORONTO  
SCHOOL OF CONTINUING STUDIES

## Module 4 – Section 3

# Modeling the Flow of Time

# Agent Percepts Over Time

- Now we have a problem: which percepts' occurrence are dependent on time? (why not just index by the location of the agent when the percept occurred?)
- So we need a way of indexing percepts by time
- We have the same need for the agent's state (location, pose, whether it still has the arrow) and the status of the gold and Wumpus
- We call an aspect of the world that changes over time a **fluent**
- It would be convenient to create a **transition model** to capture how fluents change between timesteps as a result of actions taken by the agent

# Percepts and the Transition Model

$agentLocation_{x,y}^t \Rightarrow (breeze^t \Leftrightarrow breezePerceptible_{x,y})$   
 $agentLocation_{x,y}^t \Rightarrow (stench^t \Leftrightarrow stenchPerceptible_{x,y})$

Transition Model:

$agentLocation_{1,1}^0 \wedge facingEast^0 \wedge forward^0$   
 $\Rightarrow (agentLocation_{2,1}^1 \wedge \neg agentLocation_{1,1}^1).$

...

Question: How many of these **effects axioms** would we need?

# Successor-State Axioms

Better to describe each fluent in terms of how it evolves over time:

$$f^{t+1} \Leftrightarrow \text{actionCauses}F^t \vee (f^t \wedge \neg \text{actionCauses} \neg F^t)$$

e.g.

$$\text{haveArrow}^{t+1} \Leftrightarrow (\text{haveArrow}^t \wedge \neg \text{shoot}^t)$$

$$\text{agentLocation}_{1,1}^{t+1}$$

$$\Leftrightarrow (\text{agentLocation}_{1,1}^t \wedge (\neg \text{forward}^t \vee \text{bump}^{t+1}))$$

$$\vee (\text{agentLocation}_{1,2}^t \wedge (\text{south}^t \wedge \text{forward}^t))$$

$$\vee (\text{agentLocation}_{2,1}^t \wedge (\text{west}^t \wedge \text{forward}^t)).$$



UNIVERSITY OF TORONTO  
SCHOOL OF CONTINUING STUDIES

## Module 4 – Section 4

# Planning

# Planning

- So far, our agent has used domain-independent heuristics that would allow it to generate a plan but it would need to handle a potentially huge space of states and actions
- For example, for WumpusWorld:
  - 4 agent orientations x max # of time steps x 16 possible locations
- We can deal better with this space if we factor it
- Our approach: Represent each state as a conjunction of ground, functionless fluents that describe its features

# Assumptions for the Planning Method to Follow

- Single, deterministic agent that predicts the consequences of its actions
- No events outside its control
- The environment is fully observable
- Time progresses discretely from one state to the next
- Goals are predicates of states to be achieved

# Representation Using Planning Languages

- Current State
- Action
  - Precondition
  - Effect
- Resulting State

*goal( $\neg firePresent$ )*

*action(activate),*

*PRECOND: at(station)  $\wedge \neg batteryLow \wedge firePresent$*

*EFFECT:  $\neg at(station) \wedge at(enroute)$*

# Forward Planning

- Create a graph of the relevant space
- Use search such as A\*
- Optimal planning is PSPACE which is worse than NP but often finding a non-optimal solution can be P if we have good heuristics
- Common heuristics:
  - Ignore preconditions
  - Ignore delete lists
  - State abstraction
  - Decomposition

# Planning With Multiple Agents

- Multi-agent vs. multi-body
- Coordination and communication
- Hierarchical control
- Team play
  - Cohesion
  - Separation
  - Alignment



UNIVERSITY OF TORONTO  
SCHOOL OF CONTINUING STUDIES

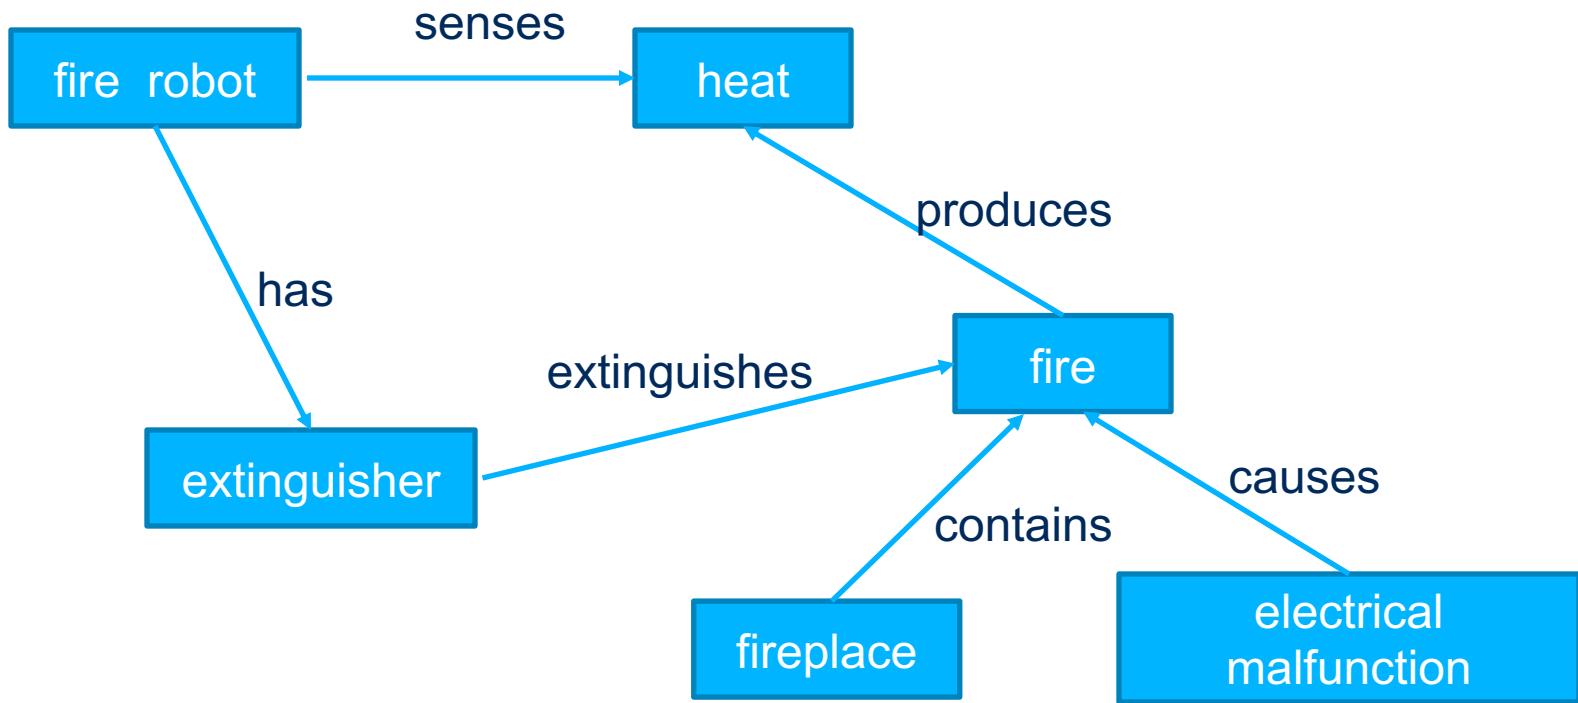
## Module 4 – Section 5

# Knowledge Representation

# Ontologies

- An ontology is a specification of the meanings of the symbols in a knowledge base or information system
- We want to model
  - Things
  - Relationships
  - Properties of things
- A popular form is to model relationships as a **triple**:  
 $prop(Individual, Property, Value)$
- Triples correspond to *subject verb object* form and can also be written  $prop(subject, verb, object)$  or  $verb(subject, object)$
- Triples can be interpreted as forming a directed graph called a **semantic network** or **knowledge graph**

# An Example Semantic Network

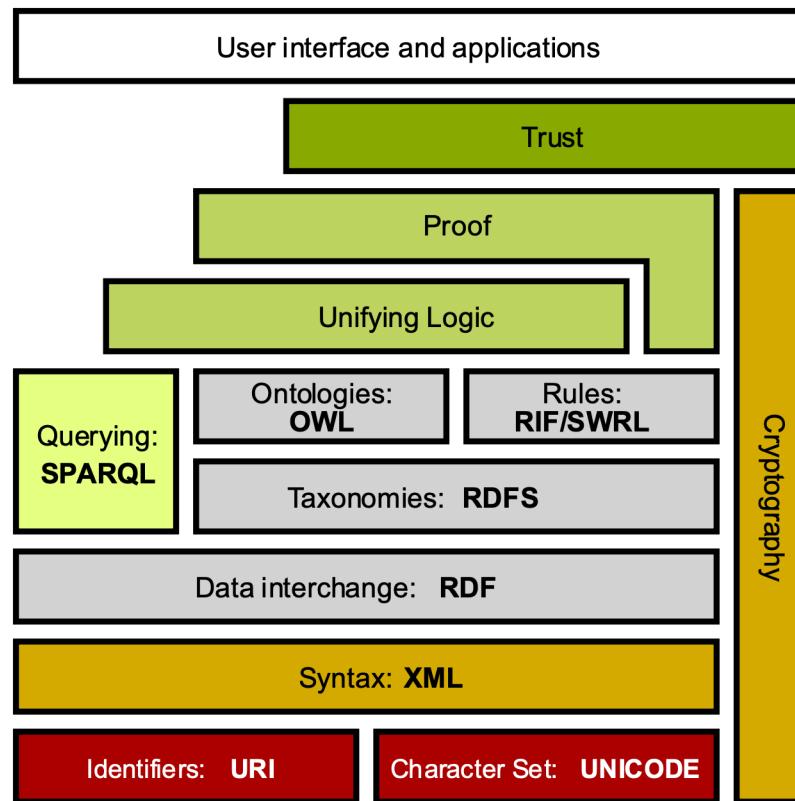


# Classes and Part-of

- It is a natural extension to recognize groups of things that are similar
- We can introduce a distinguished Property **is\_type\_of**
- Similarly, we may distinguish **is\_part\_of**

# Semantic Web

- Berners-Lee's vision of the WWW is a world-wide semantic network
- Based on triples and semantic markup





UNIVERSITY OF TORONTO  
SCHOOL OF CONTINUING STUDIES

## Module 4 – Section 6

# Resources and Wrap-up

# Resources

- Intro to Prolog: <https://metalevel.at/prolog>
- Bratko, Ivan. Prolog Programming for Artificial Intelligence. 4<sup>th</sup> Ed. 2012.
- SWI-Prolog: <https://www.swi-prolog.org/>
- Gnu Prolog: <http://www.gprolog.org/>
- Lambda Prolog:  
<http://www.lix.polytechnique.fr/Labo/Dale.Miller/IProlog/>
- [https://en.wikipedia.org/wiki/DPLL\\_algorithm](https://en.wikipedia.org/wiki/DPLL_algorithm)
- Scryer Prolog: <https://github.com/mthom/scryer-prolog>
- Datalog vs. Prolog:  
<http://www.cs.tau.ac.il/~msagiv/courses/pl14/prolog.pdf>

# Summary

- Logic programming was “AI” until Deep Learning appeared
- Deep Learning is good at pattern recognition and soft logic but not reasoning\*
- Logic programming enables declarative problem-solving
- Planning can be expressed as a search problem and solved using logic programming but is computationally a “Hard” problem
- Knowledge can be represented with very simple constructs (e.g. triples) and given meaning by ontologies

\* yet?

# Next Week

- Logic programming treats propositions as true/false dichotomies: *black and white thinking*...
- We live in an uncertain world, so: we'll use Bayesian techniques to enable probabilistic inference



UNIVERSITY OF TORONTO  
SCHOOL OF CONTINUING STUDIES

# Any questions?



# Thank You

Thank you for choosing the University of Toronto  
School of Continuing Studies